
Pre-Reads

The AI Director

Visual Storytelling with AI

 Reading time: **~15 minutes**

 What you'll need: **Just a laptop and internet for the live session**

What to Expect From the Session

You'll watch real visual content being created live using AI — from a text description to finished, professional-quality output. You'll see how different AI tools handle different types of visual tasks, and why choosing the right tool for the right job matters.

By the end of the session, you'll understand:

- How AI generates images and videos from text descriptions

- Why prompting for visuals is completely different from prompting for text
- A framework for writing prompts that give you director-level control over AI visuals
- How to think about which AI tool to use for which type of visual task
- What's realistic to expect from AI video today — and where the limits are

✓ **This is a watch-and-learn session.** Practice materials will be shared afterward.

The Big Idea Behind This Session

When most people try AI image or video tools, they type something like "a beautiful sunset on a beach" and get a generic result. They might think, "That's cool, but I could find a better stock photo." And they'd be right.

The difference between someone who gets generic outputs and someone who gets stunning, professional results isn't the tool. It's the **way they describe what they want**.

Think about it this way. If you hired a photographer and said "take a nice photo," you'd get something random. But if you said "shoot from a low angle, golden hour lighting, shallow depth of field, subject walking toward camera on a cobblestone street" — now you're directing. You're giving specific creative instructions that shape the final result.

That's the shift this session teaches. You're not just "using an AI tool." You're learning to **direct AI** the way a filmmaker directs a crew. The AI is your camera operator, your lighting designer, and your editor — but you're the director. And the quality of the final product depends entirely on how well you communicate your vision.

Key Terms You'll Hear in the Session

Don't memorize these. Read through them once so they sound familiar when they come up.

Term	What It Means (Plain English)
------	-------------------------------

Text-to-Image	You type a description, AI generates an image. Think of it like describing a scene to a painter. The more specific your description, the closer the image matches what you imagined. Vague descriptions get vague results.
Text-to-Video	Same idea, but the output is a moving video clip instead of a still image. You're describing what happens, how the camera moves, and what the mood feels like — not just what things look like.
Image-to-Video	You give AI a still image, and it brings it to life with motion. Like those "living photos" from Harry Potter — a frozen moment that starts to move. Useful when you've already got a perfect image and want to animate it into a video clip.
Diffusion Model	The technology behind most AI image and video tools. Think of it like a sculptor who starts with a rough, noisy block of marble and gradually chips away until a clear image emerges. The AI starts with visual noise and step by step refines it into a coherent image based on your description.
Upscaling	AI-generated video often comes out a bit soft or blurry. Upscaling sharpens details and increases resolution — like enhancing a phone photo to look like it was shot on a professional camera. It's the finishing step that makes the difference between "that looks like AI" and "I can't tell that's AI."
Prompting for Visuals	When you prompt ChatGPT or Claude, you're writing plain English. Visual prompting adds a new vocabulary — camera angles, lighting styles, lens types, colour palettes. The session will introduce the key terms and show you how to use them.

The Director's Mindset

This is the most important concept in the session. When you prompt AI for text, you think about **what you want it to say**. When you prompt AI for visuals, you need to think about **what you want it to show** — and how.

A film director doesn't just say "film a coffee shop scene." They make dozens of decisions:

What's in the frame? A person? A product? A landscape?

What's happening? Someone walking? Steam rising? A hand reaching?

Where is the camera? Close-up? Wide shot? Overhead? Low angle?

What's the lighting? Bright and cheerful? Dark and moody? Golden hour?

What's the mood? Energetic? Calm? Dramatic? Cosy?

These are the same decisions you'll learn to make when prompting AI visual tools. The session introduces a structured framework that breaks this down into simple, repeatable steps. No photography or filmmaking background needed.

The key insight: **the more precisely you can describe what you see in your mind, the more precisely AI can create it.** The same principle from text prompting applies here — specificity wins — but the vocabulary is different.

The AI Visual Tool Landscape

There are dozens of AI image and video tools available right now, and new ones launch regularly. The session will walk through the major players, but here's what matters more than any specific tool name:

Different tools are better at different things. Some excel at generating still images. Some are best at short video clips. Some handle camera movement well. Some can generate audio alongside video. Some are good at keeping a character looking consistent across multiple shots.

The skill isn't memorizing which tool does what — because that changes every few months. The skill is understanding what kind of output you need and knowing how to evaluate which tool fits that need. That's a transferable skill that stays useful no matter how fast the tools evolve.

You'll also learn about the difference between:

- **Free tiers vs paid plans** — Most tools offer free trials or limited free usage. Paid plans typically range from \$20–50 per month. The session will be transparent about what things cost.

- **Running one tool vs combining multiple tools** — Just like the workflows session taught you to chain AI text tools together, this session shows how combining visual tools produces results that no single tool could achieve alone.

Come Prepared: One Visual Idea

During the session, you'll see how detailed visual prompts produce dramatically better results than vague ones. To get the most out of it, come with a visual idea in mind.

Think of something like:

- A product you'd want to create an ad for
- A social media post you'd want eye-catching visuals for
- A short video clip for your brand, business, or personal project
- A thumbnail or banner image for content you create

It doesn't have to be a fully formed concept. Just have something visual in mind so that when you see the techniques demonstrated, you can immediately think about how you'd apply them to your own work.

A Few Things to Keep in Mind

No design or filmmaking experience needed. Everything is explained from scratch. If you've never touched a video editing tool, that's fine.

Watch the thinking, not the specific tools. Tool names and features change constantly. The director's mindset — knowing how to describe what you want — is what lasts.

AI visuals are improving fast. What you see in the session represents where the technology is right now. Six months from now, it'll be even better. The frameworks you learn will still apply.

The results might surprise you. The gap between "AI-generated video" and "professional production" is closing faster than most people realise. Come with an open mind about what's possible.

See you in class!