

From Prompts to AI Workflows

Learn one simple workflow you can use tomorrow — then see where it can go.

 **Session led by** — Dileep Karri

 **Goal** — one reliable workflow you can use immediately, plus the advanced version for high-stakes work

 **For** — first-time and early AI workflow users

HOW TO READ THESE NOTES

The Thesis of the Session

These notes follow the session Dileep Karri led on moving from one-off prompts to repeatable AI workflows. They are written so you can rebuild everything yourself: every prompt is reproduced verbatim, every comparison is kept, and the analogies used in the room are included because they are the part that makes the ideas stick.

One sentence carries the whole session: **you do not need to become an AI expert — you need one good prompt and the judgment to know when to go deeper.** Everything below is in service of that.

PART 1

You Already Know How to Prompt. Now Learn How to Think.

A single prompt gives you an answer. A workflow gives you a **reliable outcome**. That is the entire shift this session is about — and it is a shift in thinking, not in tools.

Random prompting	Workflow approach
“Analyze this data”	A precise question, a checked output, a clear action
A different result every time	Repeatable quality
No way to improve it	Clear steps to refine
Hard to explain to others	Easy to hand off and teach

THE FRAMING

A single prompt is like asking a stranger one question and walking away. A workflow is like having a process — same input quality, same output quality, every time.

The problem with “Analyze this data”

You upload a spreadsheet and type “analyze this and tell me what you find.” Back comes a generic chart, an obvious observation, and no decision. The reason is simple: you never told the AI what decision you are trying to make — and you might not yet know how to frame that yourself.

The fix is **not** a harder prompt. It is making the AI do the hard part: asking *you* the right questions first.

THE DOCTOR ANALOGY

Walking up and saying “analyze this data” is like walking into a doctor's office and saying “fix me.” A bad doctor guesses. A good doctor takes your history first. The skill isn't asking better questions — it's hiring something that asks them for you.

PART 2

The Analyst Prompt — Your Everyday Workflow

Instead of you framing the perfect question, you paste **one prompt** that turns Claude into a senior analyst. It does what a good consultant does: it interviews you first, waits for your answers, and only then delivers a structured, executive-ready analysis. You do not need to know the framework — the prompt carries it for you.

WHY THIS MATTERS

The difference between a junior hire who runs off and does something the moment you ask, and a consultant who first sits you down and asks “what decision are you actually trying to make?” The second is worth ten times more. This one prompt rents you that second person. “Simple” doesn't mean “weak” — it means the effort is hidden inside the prompt, not loaded onto you.

The Analyst Prompt — copy and use

PASTE THIS INTO CLAUDE WITH YOUR DATA

You are a senior McKinsey-style analyst. I am going to give you some data (or describe a decision). Do NOT analyze anything yet.

FIRST, interview me. Ask me up to 6 sharp clarifying questions, in one batch, to understand:

- The decision I'm actually trying to make
- What a good outcome looks like (the metric that matters)
- The scope (time period, segment, geography)
- What I already know and what data I have
- Any constraints or assumptions you should respect

WAIT for my answers.

THEN deliver your analysis in this structure:

1. Executive Summary - the answer and why, in 5 lines
2. What the Data Shows - key findings with the numbers behind them
3. What This Means - implications and risks
4. Recommendation - one clear action, expected impact, main risk, first step
5. What I'm Not Sure About - weakest assumptions and what would change the answer

Keep it crisp. No filler. Use plain language a busy executive understands.

The two power lines

“Do NOT analyze anything yet” and “WAIT for my answers.” Those two sentences are what stop the AI from running off. You paste this prompt once and never write it again.

Why it works

Without the Analyst Prompt	With it
You must know what to ask	The AI asks you
You get a generic data dump	You get a decision
A different structure every time	The same five-part structure every time
Easy to fool yourself	“What I'm Not Sure About” forces honesty

The trick is that you moved the hard thinking from *you* to *the prompt*.

TWO ANALOGIES THAT LANDED

“Analyze this data” is throwing random ingredients in a pan and hoping. The Analyst Prompt is a recipe card — same dish, every cook, every time. And the last section, “What I’m Not Sure About,” is the AI’s confession box: most people never ask AI what it’s unsure about, so it never tells them. This prompt makes it mandatory.

The live demo — what to watch for

The demo needs no setup, no plugins — just Claude and a file:

1. Upload a real dataset (CSV/Excel) straight into Claude.
2. Paste the Analyst Prompt.
3. Answer the clarifying questions Claude asks back — even imperfectly; it still works.
4. Read the structured analysis it produces, and scroll to “What I’m Not Sure About.”

The moment that earns trust is when Claude fires its questions back at you — it is interviewing you, exactly as the prompt promised — and again when it openly states its weakest assumptions.

PART 3

ACR — The Advanced Workflow

The Analyst Prompt is your everyday tool. ACR is what you *graduate into* — not where you start. Reach for it when the decision involves serious money, people, or deadlines; when you need to audit every step rather than trust one output; or when you will run the process repeatedly and want it fixed and teachable.

AUTOMATIC VS MANUAL

The Analyst Prompt is an automatic car — it gets you anywhere with no thinking about gears. ACR is manual transmission — more effort, but full control when the road gets serious. You do not need ACR for most things.

The three steps

Three deliberate steps, and you drive each one:

- **A = Ask** — frame the question yourself: decision, metric, scope.
- **C = Check** — make the AI argue against its own answer.

- **R = Recommend** — force the insight into one executive action.

Where the Analyst Prompt does this *for* you, ACR makes you do it *by hand* — which is exactly the point when the stakes are high and someone will challenge your numbers in a room.

A = Ask

You frame the question with the decision, metric and scope named explicitly.

THE ASK TEMPLATE

You are a business analyst helping me make a specific decision.

DECISION: [What action am I choosing between?]

METRIC: [What number tells me if this worked?]

SCOPE: [Time period + segment + geography]

CONTEXT: [What do I already know? What data do I have?]

Return:

1. A single, precise analysis question
2. Three sub-questions that support it
3. The data fields I need
4. What a strong answer looks like vs a weak one

Weak ask: “Analyze this dataset.” **Strong ask:** names the decision, the metric, and the scope. A fully worked example used in the room:

WORKED EXAMPLE — A REAL SALES DECISION

You are a business analyst helping me make a specific sales decision.

I have an online sales dataset with 240 transactions from Jan 1, 2024 to Aug 27, 2024.

The dataset fields are: Transaction ID, Date, Product Category, Product Name, Units Sold, Unit Price, Total Revenue, Region, Payment Method.

Product categories: Electronics, Home Appliances, Clothing, Books, Beauty Products, Sports.

Regions: North America, Europe, Asia.

DECISION:

Which product category and region should we prioritize for growth in the next quarter?

METRIC:

Total Revenue, supported by Units Sold and average revenue per transaction.

SCOPE:

Jan 1, 2024 to Aug 27, 2024, across all regions, categories, and payment methods.

CONTEXT:

We want to identify where the strongest revenue opportunities are - which categories generate the most revenue, which regions are strongest, and whether any category-region combinations look especially promising.

Return:

1. A single, precise analysis question
2. Three sub-questions that support it
3. The data fields needed to answer each sub-question
4. What a strong answer would look like vs a weak answer

THE RESTAURANT ANALOGY

“Bring me food” versus “I want the grilled fish, no butter, lemon on the side, in 20 minutes.” Same kitchen, wildly different result. With the Analyst Prompt the AI asked these questions for you; here, you write them — more work, more precision, more control.

C = Check

After the analysis, you make the AI proofread itself.

THE CHECK TEMPLATE

Review your previous analysis critically.

1. What are the 3 weakest assumptions you made?
2. What data would change your conclusion if it were different?
3. Rate your confidence: High / Medium / Low for each insight
4. What did you NOT consider that might matter?
5. If someone argued against your conclusion, what's their strongest point?

If it rates everything “High confidence,” push back: *“Be more honest. What are you least sure about and why?”*

THE PROOFREADING ANALOGY

You don't trust the first draft of an email to your boss — you reread it with a critic's eye. The Check step makes the AI proofread itself. You're asking it to argue against itself, and it's

surprisingly good at it, if you make it.

R = Recommend

Analysis without a recommendation is just trivia.

THE RECOMMEND TEMPLATE

Based on the analysis and the self-critique, create an executive recommendation.

RECOMMENDED ACTION: [One sentence. What should we do?]
 WHY NOW: [2-3 sentences. Why this, why now?]
 EXPECTED IMPACT: [Metric + estimated range]
 MAIN RISK: [What could go wrong + mitigation]
 NEXT STEP: [One concrete action for the next 48 hours]

Under 200 words. No jargon. Understandable in 30 seconds.

Every recommendation needs four parts: **Action, Impact, Risk, Next step.**

WEATHER REPORT VS UMBRELLA

“70% chance of rain” is analysis. “Take the umbrella, leave by 8” is a recommendation. Most AI output stops at the weather report — R forces it to hand you the umbrella. The bar: could you forward it to your manager right now?

The full cycle — and the real skill

Run **Ask** → **Check** → **Recommend** end to end: upload a dataset, frame the Ask, run the Check, produce the Recommendation. Both the Analyst Prompt and ACR are valid — choosing the right one is the actual skill. *Simple prompt for most days. ACR for the day it really matters.*

PART 4

The AI Multi-Tool Workflow

So far everything lived inside one AI. But no single AI tool is the best at everything. The magic happens when you combine the right tools for the right jobs — like assembling a

team of specialists instead of relying on one generalist. You wouldn't hire one person to design, research, and analyze; you'd hire a team.

What you'll build: a research-backed strategy output (a slide deck, report, or plan) grounded in real data, styled in your brand's look, and structured the way a consultant would deliver it. The same framework answers “what product to launch,” “which market to enter,” “how to position against competitors,” or “what content strategy to follow.”

The framework: 3 roles, 3 tools

Role	Tool	What it does
 The Designer	Gemini	Captures your brand's look and feel
 The Researcher	Perplexity / ChatGPT	Browses the internet, finds fresh data
 The Analyst	NotebookLM	Synthesizes everything, gives structured answers



WHY THE ORDER IS FIXED

You decorate the house (brand identity) before you decide what furniture to buy (research) before you arrange the room (analysis). Out of order means redoing work. Say it back: “Designer, Researcher, Analyst – in that order, every time.”

The six steps

Step 1 – Capture your brand identity

Tool · Gemini (gemini.google.com)

Share your website URL and ask Gemini to read it and produce a reusable style guide:

"Look at this website - [URL]. Analyse the brand's design language, colour palette, typography, tone of voice, and overall personality. Summarise it as a style guide I can reuse."

Save the output — you need it in Step 5. This is your brand's fingerprint; skip it and every output looks like a stranger made it. Thirty seconds now saves an hour of “this doesn't feel like us” later.

Step 2 — Gather research

Tool · Perplexity.ai (or ChatGPT with browsing, or Gemini Deep Research)

"I run [your business - URL]. Research my competitors, market trends, pricing, and customer pain points. Summarise with sources."

Copy and save the full output — text and source links. Why separately? NotebookLM cannot browse the internet; you bring the research to it.

THE LIBRARY ANALOGY

NotebookLM is a brilliant analyst locked in a library with no internet — it can only reason about what you carry into the room. Perplexity is the runner who goes outside and brings back today's newspaper. Don't skip the runner.

Step 3 — Load everything into NotebookLM

Tool · NotebookLM (notebooklm.google.com)

Create a notebook and add as sources:

- Your company website URL
- The Perplexity research (paste as a Note, then convert to a Source)
- The brand style guide from Step 1 (same: Note → Source)
- Any extra material — competitor sites, reports, reviews

Think of each source as adding another expert to the room. A panel of one is an opinion; a panel of six with citations is a briefing. Quality over quantity, though — junk sources add a loud, wrong voice.

Step 4 — Configure the notebook voice

⚠ This is the step most people skip. Don't.

This is the single highest-leverage step. It is the briefing you would give a new hire on day one. Skip it and you get a generic AI; do it and you get *your* consultant.

Open **Settings / Configure** and paste this, replacing the bracketed parts:

NOTEBOOKLM CONFIGURATION PROMPT

ROLE

You are a strategy consultant for [YOUR BRAND NAME],
a [one-line brand description, e.g. "premium kitchen appliance brand in India"].
You are structured, data-driven, and give clear recommendations.

RULES

- Only use the sources provided. Do not make things up.
- Use numbers and citations wherever possible.
- Keep language crisp. Short paragraphs, clear headings, no filler.
- If information is missing, say so and suggest what data to collect.
- When making slides, follow the brand design language from the sources.

WHEN I ASK FOR AN ANALYSIS, STRUCTURE IT AS:

1. Executive Summary - what to do and why (10 lines max)
2. Market Map - how the market is segmented and where the money is
3. Competitor Landscape - who competes and how they're positioned
4. Customer Insights - what customers want and what frustrates them
5. Recommendations - 2-3 options ranked by priority with reasons
6. Scoring Table - rate each option (1-5): market size, brand fit, differentiation, profitability, speed to launch
7. Open Questions - what still needs answers before deciding

Step 5 — Ask your questions

Ask in this order — big picture first, then narrow:

#	Ask this	You get
1	<i>“Give me a full analysis.”</i>	The structured report from your prompt
2	<i>“What should we do next? Rank options in a table.”</i>	Prioritised recommendations
3	<i>“Create a slide deck on this for my CEO. Use the brand design language.”</i>	Strategy deck in your brand's style
4	<i>“What are the biggest risks of [Option X]?”</i>	Risk analysis with mitigations
5	<i>“Summarise what customers are saying.”</i>	Voice-of-customer summary

Deck format tip: ask for a “detailed deck” for internal teams (more data), or “presenter slides” for live talks (fewer slides, bigger points). Asking “what are the risks” before “what are the options” is like asking about side effects before you've picked the medicine.

Step 6 — Export and share

- **Slide deck?** Copy into Google Slides or PowerPoint, then apply the Step 1 brand guide.
- **Document?** Copy into Google Docs and share.
- **Need more depth?** Add more sources anytime — the notebook gets smarter with every source.

The output isn't trapped; it travels into the tools you already use. And the loop is reusable: next quarter, same notebook, add fresh research, ask again. You built an asset, not a one-off.

Why the multi-tool approach works

Problem with one tool	How multi-tool solves it
AI makes things up	NotebookLM only uses YOUR sources — near-zero hallucination
Output looks generic	Gemini captures your real brand identity first
Research is shallow	Perplexity browses the live internet for fresh data
Analysis is unstructured	The configured prompt forces consultant-level structure

Each tool covers another tool's blind spot. That is the entire reason multi-tool beats single-tool.

Adapt this to any problem

Your problem	Designer step	Researcher step	Analyst step
Product to launch?	Brand style guide	Competitor & market research	Launch recommendations
Content to create?	Brand tone of voice	Trending topics & audience data	Content calendar
Enter a new market?	Brand positioning	Market size & regulations	Entry strategy
Improve pricing?	Current pricing page	Competitor pricing & reviews	Pricing recommendations

One rule to remember

“Designer, Researcher, Analyst — in that order, every time.” 🎨 → 🔍 → 🧠. Total time: roughly 30–45 minutes for a strategy deck that would take days manually. That is the line people repeat to a colleague.

PART 5

From Copy-Paste to Connected: Connectors

Right now your AI lives in a box. To use it with real work you become the courier: open Drive, find the file, download it; open the AI, upload it, get the analysis; open Slides, rebuild it by hand; copy-paste back and forth between tabs. You are the middleman. **Connectors remove the middleman.**

A connector lets the AI reach into an app directly — read your Drive, draft in Gmail, build in Canva — instead of you shuttling files around. It is one standard, originally proposed by Anthropic (makers of Claude), now adopted by ChatGPT, Gemini and others. **Connect once, works everywhere** — like a universal adapter for AI.

THE UNIVERSAL ADAPTER

Before, every country (app) needed its own custom plug — exhausting, breaks constantly. Connectors are the one adapter that fits every socket. It's plumbing — you don't need to understand plumbing to turn on a tap.

Setting it up

On Claude	On ChatGPT (paid plan)
Open Claude.ai	Open ChatGPT
Click Tools → Add Connectors	Settings → Connectors (or Add Sources in a chat)
Browse connectors (Drive, Notion, Canva, Gmail...)	Pick the app (Drive, SharePoint, GitHub...)
Click Connect on the app you want	Click Connect and authenticate
Authenticate (log in, approve access)	Reference the connected source in your prompt
Done — Claude auto-detects which connector to use	Same idea, two platforms

The magic moment is that you don't say “use the Drive connector” — you just say “*summarize my latest Drive file*” and it figures out which key to use. A good assistant doesn't ask “which filing cabinet?” — they just know.

What you can connect

Common connectors: Google Drive, Gmail, Canva, Notion, Slack, GitHub, Stripe, and more.

	Claude	ChatGPT
Free tier	Yes (limited)	No — paid only
Paid (~\$20/mo)	More usage + connectors	Connectors included
Auto-detects connector	Yes	Partial

The honest framing: if your company already pays for one of these, you already have this — most people just never turned it on. In a live demo, the point is the absence of work — no download, no upload, no copy-paste. That is three steps gone from every task you do.

⚠ Safety & privacy — three rules before you connect anything

- **Allow Once, not Always.** Start with the one-time option; widen access only after you trust it. It's a houseguest's day pass, not a permanent key to your home.
- **Verify before you act.** Read AI-written emails and messages before they send. Never auto-send.
- **Mind the tier.** Free plans may use your data for training. Don't connect confidential company data on free tiers; use paid or enterprise for sensitive work.

PART 6

Pulling It Together

Common mistakes and fixes

Mistake	Fix
Starting with “analyze this data”	Use the Analyst Prompt — let the AI interview you first
Trusting the first AI output	Make it critique itself (the Check step)
Analysis with no action	End every analysis with a recommendation
Skipping the NotebookLM config step	Always brief your AI consultant before asking
Connecting 10 apps at once	Start with one connector; add more later
Sharing sensitive data on free plans	Check the plan's privacy policy first

If you fix only one habit, replace “analyze this data” with the Analyst Prompt.

Workflows vs agents

	Workflow (today)	Agent (what comes next)
You control	Every step	The goal only
Path	Predefined by you	AI decides dynamically
Best for	Repeatable tasks	Ambiguous, multi-step problems
Risk	Low — you check everything	Higher — AI makes more decisions
Start here?	Yes	Not yet

Master workflows first. Agents come after you trust the process — you learn the recipe before you trust the chef.

KEEP THIS

Quick Reference Card

Everyday — The Analyst Prompt

```
You are a senior McKinsey-style analyst. Do NOT analyze yet.  
FIRST interview me (up to 6 questions): decision, success metric,  
scope, what I know/have, constraints. WAIT for answers.  
THEN: 1) Exec Summary 2) What the data shows 3) What it means  
4) Recommendation (action, impact, risk, next step)  
5) What I'm not sure about. Crisp, plain language.
```

Advanced — ACR

```
ASK      -> DECISION / METRIC / SCOPE / CONTEXT  
CHECK    -> 3 weakest assumptions? what would change it?  
          confidence per insight? what's missing? counter-argument?  
RECOMMEND -> ACTION / WHY NOW / IMPACT / RISK / NEXT STEP
```

The whole session in one line, from Dileep: you don't need to become an AI expert — you need one good prompt and the judgment to know when to go deeper. Simple prompt for most days. ACR for the day it really matters.